JiZhiVine: A Decision-only Foundation Model for Personalized Ranking and Recommendation

Yi Miao, Bencheng Yan, Jiahui Zhao, Jiawei Chen, Chengcheng Guo, Haoran Zhang, Haiyang Hu, Yuanzhong Liu, Yuanzhong Xu, Yanqing Hu

Alibaba Group
{yi.miao,bencheng.yan,jiahui.zhao,jiawei.chen.zju,chengcheng.guo,haoran.zhang,haiyang.hu,yuanzhong.liu,yuanzhong.xu,yanqing.hu}@alibaba-inc.com

Abstract

Modern large recommendation systems commonly implement cascaded architectures, spanning retrieval, coarse ranking, and fine-grained ranking, to balance computational scalability with ranking accuracy. While deep learning has revolutionized each stage of this pipeline, joint optimization across the entire cascade remains an open challenge. Recent foundation model approaches show promise in unifying ranking tasks, but existing methods primarily focus on two-stage retrieval and ranking, or rely on large-scale ensemble architectures that are impractical to deploy in production. We propose JiZhiVine, a decision-only foundation model for personalized ranking and recommendation. Our model adopts a HSTU-based sequential transformer architecture with 330M parameters, trained on 1.2 trillion tokens from 400 billion user interactions spanning 6 months. We introduce a multi-stage ranking loss that jointly optimizes retrieval, coarse ranking, and fine-grained ranking stages, enabling gradients to propagate through the entire cascade. We also propose a hardware-aware mixed-precision training framework with custom CUDA kernels for efficient inference on GPU clusters. Extensive experiments on a public benchmark and our production dataset demonstrate that JiZhiVine achieves state-of-the-art performance across all ranking stages while maintaining sub-linear inference cost relative to model depth. We provide detailed ablation studies on architecture scaling, loss function design, and cascade ranking consistency, offering insights for practitioners deploying large-scale ranking systems.

1 Introduction

Figure 1: Overview of a typical recommendation cascade ranking system.

Recommendation and advertising systems serve hundreds of billions of users daily, optimizing for metrics spanning click-through rate, conversion, dwell time, and long-term user satisfaction. These systems must process massive candidate pools, from tens of millions of items at the retrieval stage to hundreds at the final ranking stage, while maintaining sub-100ms latency and high throughput. The dominant paradigm for this scale is a cascaded architecture, where retrieval, coarse ranking, and fine-grained ranking stages progressively filter and refine candidates (Figure 1).

While deep learning has revolutionized each stage of this pipeline, most optimization efforts have been local to individual stages. Joint optimization across the cascade remains challenging, as gradients do not flow between stages and each stage must be trained on its own objective. Recent foundation model approaches show promise in unifying ranking tasks, but existing methods primarily focus on two-stage retrieval and ranking, or rely on large-scale ensemble architectures that are impractical to deploy in production.

In this work, we propose JiZhiVine, a decision-only foundation model for personalized ranking and recommendation. We adopt a HSTU-based sequential transformer architecture [1] with 330M parameters, trained on 1.2 trillion tokens from 400 billion user interactions spanning 6 months. We introduce a multi-stage ranking loss that jointly optimizes retrieval, coarse ranking, and fine-grained ranking stages, enabling gradients to propagate through the entire cascade. We also propose a hardware-aware mixed-precision training framework with custom CUDA kernels for efficient inference on GPU clusters.

Extensive experiments on a public benchmark and our production dataset demonstrate that JiZhiVine achieves state-of-the-art performance across all ranking stages while maintaining sub-linear inference cost relative to model depth. We provide detailed ablation studies on architecture scaling, loss function design, and cascade ranking consistency, offering insights for practitioners deploying large-scale ranking systems.

2 Related Work

2.1 Large-scale ranking

Ranking models have evolved from gradient boosted decision trees [2] to deep neural networks. Wide & Deep [3] combines memorization and generalization by jointly training a linear model and a deep network. DeepFM [4] learns feature interactions without manual feature engineering. DCN [5] introduces cross layers that explicitly model bounded-degree feature interactions. DCN-V [6] improves the cross network with vector-wise interactions and low-rank decomposition. These models operate on pre-ranking candidates (hundreds to thousands) and typically use 2-3 layers. We show that our multi-stage ranking loss enables gradients to propagate through the entire cascade, improving consistency between retrieval and fine-grained ranking.

2.2 Sequential recommendation

Sequential recommendation models user preferences as a function of interaction history. GRU4Rec [7] applies recurrent neural networks to session-based recommendation. SASRec [8] uses self-attention to capture long-range dependencies. BERT4Rec [9] employs bidirectional attention with masked item prediction. HSTU [1] introduces hierarchical sequential transformer units with efficient attention for billion-scale industrial systems. We adopt HSTU as our backbone architecture and extend it with multi-stage ranking losses.

2.3 Cascade ranking

Cascade ranking [10] [11] [12] [13] [14] [15] [16] [17] [18] [19] [20] [21] [22] [23] [24] [25] [26] [27] [28] [29] [30] [31] [32] [33] [34] [35] [36] [37] [38] [39] [40] [41] [42] [43] [44] [45] [46] [47] [48] [49]